

Online Influence Maximization under Independent Cascade Model with Semi-Bandit Feedback

Authored by:

Branislav Kveton
Zheng Wen
Michal Valko
Sharan Vaswani

Abstract

We study the online influence maximization problem in social networks under the independent cascade model. Specifically, we aim to learn the set of "best influencers" in a social network online while repeatedly interacting with it. We address the challenges of (i) combinatorial action space, since the number of feasible influencer sets grows exponentially with the maximum number of influencers, and (ii) limited feedback, since only the influenced portion of the network is observed. Under a stochastic semi-bandit feedback, we propose and analyze IMLinUCB, a computationally efficient UCB-based algorithm. Our bounds on the cumulative regret are polynomial in all quantities of interest, achieve near-optimal dependence on the number of interactions and reflect the topology of the network and the activation probabilities of its edges, thereby giving insights on the problem complexity. To the best of our knowledge, these are the first such results. Our experiments show that in several representative graph topologies, the regret of IMLinUCB scales as suggested by our upper bounds. IMLinUCB permits linear generalization and thus is both statistically and computationally suitable for large-scale problems. Our experiments also show that IMLinUCB with linear generalization can lead to low regret in real-world online influence maximization.

1 Paper Body

Social networks are increasingly important as media for spreading information, ideas, and influence. Computational advertising studies models of information propagation or diffusion in such networks [16, 6, 10]. Viral marketing aims to use this information propagation to spread awareness about a specific product. More precisely, agents (marketers) aim to select a fixed number of influencers

(called seeds or source nodes) and provide them with free products or discounts. They expect that these users will influence their neighbours and, transitively, other users in the social network to adopt the product. This will thus result in information propagating across the network as more users adopt or become aware of the product. The marketer has a budget on the number of free products and must choose seeds in order to maximize the influence spread, which is the expected number of users that become aware of the product. This problem is referred to as influence maximization (IM) [16]. For IM, the social network is modeled as a directed graph with the nodes representing users, and the edges representing relations (e.g., friendships on Facebook, following on Twitter) between them. Each directed edge (i, j) is associated with an activation probability $w(i, j)$ that models the strength of influence that user i has on user j . We say a node j is a downstream neighbor of node i if there is a directed edge (i, j) from i to j . The IM problem has been studied under a number of 31st Conference on Neural Information Processing Systems (NIPS 2017), Long Beach, CA, USA.

diffusion models [16, 13, 23]. The best known and studied are the models in [16], and in particular the independent cascade (IC) model. In this work, we assume that the diffusion follows the IC model and describe it next. After the agent chooses a set of source nodes S , the independent cascade model defines a diffusion (influence) process: At the beginning, all nodes in S are activated (influenced); subsequently, every activated node i can activate its downstream neighbor j with probability $w(i, j)$ once, independently of the history of the process. This process runs until no activations are possible. In the IM problem, the goal of the agent is to maximize the expected number of the influenced nodes subject to a cardinality constraint on S . Finding the best set S is an NP-hard problem, but under common diffusion models including IC, it can be efficiently approximated to within a factor of $1 - 1/e$ [16]. In many social networks, however, the activation probabilities are unknown. One possibility is to learn these from past propagation data [25, 14, 24]. However in practice, such data are hard to obtain and the large number of parameters makes this learning challenging. This motivates the learning framework of IM bandits [31, 28, 29], where the agent needs to learn to choose a good set of source nodes while repeatedly interacting with the network. Depending on the feedback to the agent, the IM bandits can have (1) full-bandit feedback, where only the number of influenced nodes is observed; (2) node semi-bandit feedback, where the identity of influenced nodes is observed; or (3) edge semi-bandit feedback, where the identity of influenced edges (edges going out from influenced nodes) is observed. In this paper, we give results for the edge semi-bandit feedback model, where we observe for each influenced node, the downstream neighbors that this node influences. Such feedback is feasible to obtain in most online social networks. These networks track activities of users, for instance, when a user retweets a tweet of another user. They can thus trace the propagation (of the tweet) through the network, thereby obtaining edge semi-bandit feedback. The IM bandits problem combines two main challenges. First, the number of actions (possible sets) S grows exponentially with the cardinality constraint on S . Second, the agent can only observe the influenced portion of the network as

feedback. Although IM bandits have been studied in the past [21, 8, 31, 5, 29] (see Section 6 for an overview and comparison), there are a number of open challenges [28]. One challenge is to identify reasonable complexity metrics that depend on both the topology and activation probabilities of the network and characterize the information-theoretic complexity of the IM bandits problem. Another challenge is to develop learning algorithms such that (i) their performance scales gracefully with these metrics and (ii) are computationally efficient and can be applied to large social networks with millions of users. In this paper, we address these two challenges under the IC model with access to edge semi-bandit feedback. We refer to our model as an independent cascade semi-bandit (ICSB). We make four main contributions. First, we propose IMLinUCB, a UCB-like algorithm for ICSBs that permits linear generalization and is suitable for large-scale problems. Second, we define a new complexity metric, referred to as maximum observed relevance for ICSB, which depends on the topology of the network and is a non-decreasing function of activation probabilities. The maximum observed relevance C^* can also be upper bounded based on the network topology or the size of the network in the worst case. However, in real-world social networks, due to the relatively low activation probabilities [14], C^* attains much smaller values as compared to the worst case upper bounds. Third, we bound the cumulative regret of IMLinUCB. Our regret bounds are polynomial in all quantities of interest and have near-optimal dependence on the number of interactions. They reflect the structure and activation probabilities of the network through C^* and do not depend on inherently large quantities, such as the reciprocal of the minimum probability of being influenced (unlike [8]) and the cardinality of the action set. Finally, we evaluate IMLinUCB on several problems. Our empirical results on simple representative topologies show that the regret of IMLinUCB scales as suggested by our topology-dependent regret bounds. We also show that IMLinUCB with linear generalization can lead to low regret in real-world online influence maximization problems.

2

Influence Maximization under Independence Cascade Model

In this section, we define notation and give the formal problem statement for the IM problem under the IC model. Consider a directed graph $G = (V, E)$ with a set $V = \{1, 2, \dots, L\}$ of $L = |V|$ nodes, a set $E = \{1, 2, \dots, |E|\}$ of directed edges, and an arbitrary binary weight function $w : E \rightarrow \{0, 1\}$.

We say that a node $v \in V$ is reachable from a node $v_1 \in V$ under w if there is a directed path $p = (e_1, e_2, \dots, e_l)$ from v_1 to v in G satisfying $w(e_i) = 1$ for all $i = 1, 2, \dots, l$, where e_i is the i -th edge in p . For a given source node set $S \subseteq V$ and w , we say that node $v \in V$ is influenced if v is reachable from at least one source node in S under w ; and denote the number of influenced nodes in G by $f(S, w)$. By definition, the nodes in S are always influenced. The influence maximization (IM) problem is characterized by a triple (G, K, w) , where G is a given directed graph, $K \in \mathbb{N}$ is the cardinality of source nodes, and $w : E \rightarrow [0, 1]$ is a probability weight function mapping each edge $e \in E$ to a real number $w(e) \in [0, 1]$. The agent needs to choose a set of K source nodes $S \subseteq V$ based

on (G, K, w) . Then a random binary weight function w , which encodes the diffusion process under the IC model, is obtained by independently sampling a Bernoulli random variable $w(e) \sim \text{Bern}(w(e))$ for each edge $e \in E$. The agent's objective is to maximize the expected number of the influenced nodes: $\max_S \mathbb{E}[f(S, w)]$, where $f(S, w) = \sum_{v \in V} \mathbb{P}[v \text{ is influenced} \mid S, w]$ is the expected number of influenced nodes when the source node set is S and w is sampled according to w . It is well-known that the (offline) IM problem is NP-hard [16], but can be approximately solved by approximation/randomized algorithms [6] under the IC model. In this paper, we refer to such algorithms as oracles to distinguish them from the machine learning algorithms discussed in following sections. Let S^* be the optimal solution of this problem, and $S^\dagger = \text{ORACLE}(G, K, w)$ be the (possibly random) solution of an oracle ORACLE. For any $\epsilon \in [0, 1]$, we say that ORACLE is an (ϵ, δ) -approximation oracle for a given (G, K) if for any w , $\mathbb{P}[f(S^\dagger, w) \geq (1 - \epsilon)f(S^*, w)] \geq \delta$ with probability at least δ . Notice that this further implies that $\mathbb{E}[f(S^\dagger, w)] \geq (1 - \epsilon)\mathbb{E}[f(S^*, w)]$. We say an oracle is exact if $\epsilon = 0$.

3

Influence Maximization Semi-Bandit

In this section, we first describe the IM semi-bandit problem. Next, we state the linear generalization assumption and describe IMLinUCB, our UCB-based semi-bandit algorithm. 3.1

Protocol

The independent cascade semi-bandit (ICSB) problem is also characterized by a triple (G, K, w) , but w is unknown to the agent. The agent interacts with the independent cascade semi-bandit for n rounds. At each round $t = 1, 2, \dots, n$, the agent first chooses a source node set $S_t \subseteq V$ with cardinality K based on its prior information and past observations. Influence then diffuses from the nodes in S_t according to the IC model. Similarly to the previous section, this can be interpreted as the environment generating a binary weight function w_t by independently sampling $w_t(e) \sim \text{Bern}(w(e))$ for each $e \in E$. At round t , the agent receives the reward $f(S_t, w_t)$, that is equal to the number of nodes influenced at that round. The agent also receives edge semi-bandit feedback from the diffusion process. Specifically, for any edge $e = (u_1, u_2) \in E$, the agent observes the realization of $w_t(e)$ if and only if the start node u_1 of the directed edge e is influenced in the realization w_t . The agent's objective is to maximize the expected cumulative reward over the n steps. 3.2

Linear generalization

Since the number of edges in real-world social networks tends to be in millions or even billions, we need to exploit some generalization model across activation probabilities to develop efficient and deployable learning algorithms. In particular, we assume that there exists a linear-generalization model for the probability weight function w . That is, each edge $e \in E$ is associated with a known feature vector $x_e \in \mathbb{R}^d$ (here d is the dimension of the feature vector) and that there is an unknown coefficient vector $\theta \in \mathbb{R}^d$ such that for all $e \in E$, $w(e)$ is "well approximated" by $x_e^\top \theta$. Formally, we assume that $\epsilon = \max_{e \in E} |w(e) - x_e^\top \theta|$ is small. In Section 5.2, we see that such a linear generalization

leads to efficient learning in real-world networks. Note that all vectors in this paper are column vectors. 1 As is standard in graph theory, a directed path is a sequence of directed edges connecting a sequence of distinct nodes, under the restriction that all edges are directed in the same direction. 2 Notice that the definitions of $f(S, w)$ and $f(S, w)$ are consistent in the sense that if $w \in \{0, 1\}^E$, then $f(S, w) = f(S, w)$ with probability 1.

3

Algorithm 1 IMLinUCB: Influence Maximization Linear UCB Input: graph G , source node set cardinality K , oracle ORACLE, feature vector x_e 's, and algorithm parameters $\gamma, c \in [0, 1]$, Initialization: $B_0 = 0 \in \mathbb{R}^{d \times d}$, $M_0 = I \in \mathbb{R}^{d \times d}$ for $t = 1, 2, \dots, n$ do

1. set $\hat{w}_t \in \mathbb{R}^d$ and the UCBs as $U(e) = \text{Proj}_{[0,1]}(x_e^T \hat{w}_t + c \sqrt{\lambda_t})$ for all $e \in E$ 2. choose $S_t \sim \text{ORACLE}(G, K, U_t)$, and observe the edge-level semi-bandit feedback 3. update statistics: (a) initialize $M_t = M_{t-1}$ and $B_t = B_{t-1}$ (b) for all observed edges $e \in E$, update $M_t = M_t + x_e x_e^T$ and $B_t = B_t + x_e w_t(e)$

Similar to the existing approaches for linear bandits [1, 9], we exploit the linear generalization to develop a learning algorithm for ICSB. Without loss of generality, we assume that $\|x_e\|_2 \leq 1$ for all $e \in E$. Moreover, we use $X \in \mathbb{R}^{E \times d}$ to denote the feature matrix, i.e., the row of X associated with edge e is x_e . Note that if a learning agent does not know how to construct good features, it can always choose the naive feature matrix $X = I \in \mathbb{R}^{E \times E}$ and have no generalization model across edges. We refer to the special case $X = I \in \mathbb{R}^{E \times E}$ as the tabular case. 3.3

IMLinUCB algorithm

In this section, we propose Influence Maximization Linear UCB (IMLinUCB), detailed in Algorithm 1. Notice that IMLinUCB represents its past observations as a positive-definite matrix (Gram matrix) $M_t \in \mathbb{R}^{d \times d}$ and a vector $B_t \in \mathbb{R}^d$. Specifically, let X_t be a matrix whose rows are the feature vectors of all observed edges in t steps and Y_t be a binary column vector encoding the realizations of all observed edges in t steps. Then $M_t = I + \sum_{s=1}^t X_{ts} X_{ts}^T$ and $B_t = \sum_{s=1}^t X_{ts} Y_{ts}$. At each round t , IMLinUCB operates in three steps: First, it computes an upper confidence bound $U_t(e)$ for each edge $e \in E$. Note that $\text{Proj}_{[0,1]}(\cdot)$ projects a real number into interval $[0, 1]$ to ensure that $U_t \in [0, 1]^E$. Second, it chooses a set of source nodes based on the given ORACLE and U_t , which is also a probability-weight function. Finally, it receives the edge semi-bandit feedback and uses it to update M_t and B_t . It is worth emphasizing that IMLinUCB is computationally efficient as long as ORACLE is computationally efficient. Specifically, at each round t , the computational complexities of

both Step 1 and 3 of IMLinUCB are $O(d|E|)$. 3 It is worth pointing out that in the tabular case, IMLinUCB reduces to CUCB [7], in the sense that the confidence radii in IMLinUCB are the same as those in CUCB, up to logarithmic factors. That is, CUCB can be viewed as a special case of IMLinUCB with $X = I$. 3.4

Performance metrics

Recall that the agent's objective is to maximize the expected cumulative reward, which is equivalent to minimizing the expected cumulative regret. The cumulative regret is the loss in reward (accumulated over rounds) because of the lack of knowledge of the activation probabilities. Observe that in each round t , IMLinUCB needs to use an approximation/randomized algorithm ORACLE for solving the offline IM problem. Naturally, this can lead to $O(n)$ cumulative regret, since at each round there is a non-diminishing regret due to the approximation/randomized nature of ORACLE. To analyze the performance of IMLinUCB in such Pncases, we define a more appropriate performance metric, the scaled cumulative regret, as $R(n) = \sum_{t=1}^n \mathbb{E} [R_t]$, where n is the number of steps, $\gamma \in [0, 1]$ is the scale, and $R_t = \gamma (S_{\text{opt}} - \sum_{i=1}^t w_i f(S_i, w_i))$ is the γ -scaled realized regret R_t at round t . When $\gamma = 1$, $R(n)$ reduces to the standard expected cumulative regret $R(n)$. ³

Notice that in a practical implementation, we store M_t instead of M_t . Moreover, $M_t = M_t + \gamma \sum_{i=1}^t x_i^T e_t$

is equivalent to $M_t = M_t + \gamma \sum_{i=1}^t x_i^T e_t$
 $\sum_{i=1}^t M_t = \sum_{i=1}^t M_t + \gamma \sum_{i=1}^t x_i^T e_t$

.

4

(a)

(b)

(c)

(d)

Figure 1: a. Bar graph on 8 nodes. b. Star graph on 4 nodes. c. Ray graph on 10 nodes. d. Grid graph on 9 nodes. Each undirected edge denotes two directed edges in opposite directions.

4

Analysis

In this section, we give a regret bound for IMLinUCB for the case when $w(e) = x^T e$ for all $e \in E$, i.e., the linear generalization is perfect. Our main contribution is a regret bound that scales with a new complexity metric, maximum observed relevance, which depends on both the topology of G and the probability weight function w , and is defined in Section 4.1. We highlight this as most known results for this problem are worst case, and some of them do not depend on probability weight function at all. 4.1

Maximum observed relevance

We start by defining some terminology. For given directed graph $G = (V, E)$ and source node set $S \subseteq V$, we say an edge $e \in E$ is relevant to a node $v \in V \setminus S$ under S if there exists a path p from a source node $s \in S$ to v such that (1) $e \in p$ and (2) p does not contain another source node other than s . Notice that with a given S , whether or not a node $v \in V \setminus S$ is influenced only depends on the binary weights w on its relevant edges. For any edge $e \in E$, we define $NS_{S,e}$ as the number of nodes in $V \setminus S$ it is relevant to, and define $PS_{S,e}$ as the conditional probability that e is observed given S , $\mathbb{P}(e \in NS_{S,e} | S) = \mathbb{P}(e \text{ is relevant to } v \text{ under } S)$ and $PS_{S,e} = \mathbb{P}(e \text{ is observed} | S)$. (1) Notice that $NS_{S,e}$ only depends on the topology of G , while $PS_{S,e}$ depends on both the topology of

G and the probability weight w . The maximum observed relevance $C^?$ is defined as the maximum (over S) 2-norm of $NS_{e,e}$'s weighted by $PS_{e,e}$'s, $qP \geq 2 C^? = \max_S: \sum_{e \in E(S)} NS_{e,e}^2 PS_{e,e}$. As is detailed in the proof of Lemma 1 in Appendix A, $C^?$ arises in the step where Cauchy-Schwarz inequality is applied. Note that $C^?$ also depends on both the topology of G and the probability weight w . However, $C^?$ can be bounded from above only based on the topology of G or the size of the $qP \geq 2$ problem, i.e., $L = |V|$ and $|E|$. Specifically, by defining $CG = \max_S: \sum_{e \in E(S)} NS_{e,e}^2$, we have $C^? \leq CG = \max_S: \sum_{e \in E(S)} K^2$

$$\begin{aligned} & \leq \sum_{e \in E(S)} K^2 = O(L \sum_{e \in E(S)} 1) = O(L^2), \quad NS_{e,e} \leq 1 \\ & \text{qP} \\ & (3) \end{aligned}$$

where CG is the maximum/worst-case (over w) $C^?$ for the directed graph G , and the maximum is obtained by setting $w(e) = 1$ for all $e \in E$. Since CG is worst-case, it might be very far away from $C^?$ if the activation probabilities are small. Indeed, this is what we expect in typical realworld situations. Notice also that if $\max_{e \in E} w(e) \leq 0$, then $PS_{e,e} \leq 0$ for all $e \in E(S)$ and $PS_{e,e} = 1$ for all $e \in E(S)$, where $E(S)$ is the set of edges with start node in S , hence we have $qP \leq 0 \leq 2 C^? \leq CG = \max_S: \sum_{e \in E(S)} K^2$. In particular, if K is small, CG is much less than CG in 3

many topologies. For example, in a complete graph with $K = 1$, $CG = O(L^2)$ while $CG_0 = O(L^2)$. Finally, it is worth pointing out that there exist situations (G, w) such that $C^? = O(L^2)$. One such example is when G is a complete graph with L nodes and $w(e) = 1/(L+1)$ for all edges e in this graph. To give more intuition, in the rest of this subsection, we illustrate how CG , the worst-case $C^?$, varies with four graph topologies in Figure 1: bar, star, ray, and grid, as well as two other topologies: 5

general tree and complete graph. We fix the node set $V = \{1, 2, \dots, L\}$ for all graphs. The bar graph (Figure 1a) is a graph where nodes i and $i+1$ are connected when i is odd. The star graph (Figure 1b) is a graph where node 1 is central and all remaining nodes $i \in V \setminus \{1\}$ are connected to it. The distance between any two of these nodes is 2. The ray graph (Figure 1c) is a star graph

with $k = L/2$ arms, where node 1 is central and each arm contains either $\lfloor L/2 \rfloor/k$ or $\lceil L/2 \rceil/k$ nodes connected in a line. The distance between any two nodes in this graph is $O(L)$. The grid graph (Figure 1d) is a classical non-tree graph with $O(L)$ edges. To see how CG varies with the graph topology, we start with the simplified case when $K = |S| = 1$. In the bar graph (Figure 1a), only one edge is relevant to a node $v \in V \setminus S$ and all the other edges are not relevant to any nodes. Therefore, $CG \leq 1$. In the star graph (Figure 1b), for any s , at most one edge is relevant to at most $L/2$ nodes and the remaining edges are relevant to at most one node. In this case, $CG \leq L/2 + L = O(L)$. In the ray graph (Figure 1c), for any s , at most $O(L)$ edges are relevant to L edges are relevant to at most $O(L)$ nodes and the remaining 5 nodes. In this case, $CG = O(L^2/2 + LL) = O(L^4)$. Finally, recall that for all graphs we can bound CG by $O(L \sum_{e \in E} 1)$, regardless of K . Hence, for the grid graph

(Figure 1d) and general tree graph, $CG = O(L^2)$ since $|E| = O(L)$; for the complete graph $CG = O(L^2)$ since $|E| = O(L^2)$. Clearly, CG varies widely with the topology of the graph. The second column of Table 1 summarizes how CG varies with the above-mentioned graph topologies for general $K = |S|$.

4.2

Regret guarantees

Consider C_t defined in Section 4.1 and recall the worst-case upper bound $C_t \leq (L + K) |E|$, we have the following regret guarantees for IMLinUCB. Theorem 1 Assume that (1) $w(e) = xTe$ for all $e \in E$ and (2) ORACLE is an (γ, β) -approximation algorithm. Let D be a known upper bound on k^2 , if we apply IMLinUCB with $\gamma = 1$ and $\beta = s$

$$n|E| - c = d \log(1 + 2 \log(n(L + 1 + K))) + D, \quad (4)$$

then we have

$$n|E| - c \leq d \log(1 + 2 \log(n(L + 1 + K))) + D, \quad (4)$$

$$n|E| - c \leq d \log(1 + 2 \log(n(L + 1 + K))) + D, \quad (4)$$

$$n|E| - c \leq d \log(1 + 2 \log(n(L + 1 + K))) + D, \quad (4)$$

Moreover, if the feature matrix $X = I$ (i.e., the tabular case), we have

$$n|E| - c \leq d \log(1 + 2 \log(n(L + 1 + K))) + D, \quad (4)$$

$$n|E| - c \leq d \log(1 + 2 \log(n(L + 1 + K))) + D, \quad (4)$$

$$n|E| - c \leq d \log(1 + 2 \log(n(L + 1 + K))) + D, \quad (4)$$

Please refer to Appendix A for the proof of Theorem 1, that we outline in Section 4.3. We now briefly comment on the regret bounds in Theorem 1. Topology-dependent bounds: Since C_t is topology-dependent, the regret bounds in Equations 5 and 7 are also topology-dependent. Table 1 summarizes the regret bounds for each topology discussed in Section 4.1. Since the regret bounds in Table 1 are the worst-case regret bounds for a given topology, more general topologies have larger regret bounds. For instance, the regret bounds for tree are larger than their counterparts for star and ray, since star and ray are special trees. The grid and tree can also be viewed as special complete graphs by setting $w(e) = 0$ for some $e \in E$, hence complete graph has larger regret bounds. Again, in practice we expect C_t to be far smaller due to activation probabilities.

The regret bound for bar graph is based on Theorem 2 in the appendix, which is a stronger version of Theorem 1 for disconnected graph.

6

ray graph

$$CG \text{ (worst-case } C_t) = O(K) = O(LK) = O(L^2K)$$

tree graph

$$O(L^2)$$

grid graph

$$O(L^2)$$

complete graph

$$O(L^2)$$

topology	bar	graph	star	graph
3				
3				
$R_{??}(n)$	for general X	$e(dK \sqrt{n}/(??))$	O	
$? e dL$	$32 Kn/(??)$	O		
$? e dL$	$74 Kn/(??)$	O		
$e dL2$	$\sqrt{n}/(??)$	O		
$e dL2$	$\sqrt{n}/(??)$	O		
$e dL3$	$\sqrt{n}/(??)$	O		
$R_{??}(n)$	for $X = I$	$?$		
$e L$	$Kn/(??)$	O	$?$	
$e L2$	$Kn/(??)$	O		
$? e L$	$94 Kn/(??)$	O		
$e L$	$52 \sqrt{n}/(??)$	O		
$e L$	$52 \sqrt{n}/(??)$	O		
$e L4$	$\sqrt{n}/(??)$	O		

Table 1: CG and worst-case regret bounds for different graph topologies.

Tighter bounds in tabular case and under exact oracle: Notice that for the tabular case with $e = E$ tighter regret bounds are obtained in Equations 7 and 8. feature matrix $X = I$ and $d = E$, $O(e)$ Also notice that the $O(1/(??))$ factor is due to the fact that ORACLE is an $(?, ?)$ -approximation oracle. If ORACLE solves the IM problem exactly (i.e., $? = ? = 1$), then $R_{??}(n) = R(n)$. Tightness of our regret bounds: First, note that our regret bound in the bar case with $K = 1$ matches the regret bound of the classic LinUCB algorithm. Specifically, with perfect linear generalization, this case is equivalent to a linear bandit problem with L arms and feature dimension d . From Table 1, $e(d\sqrt{n})$, which matches the known regret bound of LinUCB that can our regret bound in this case is O be obtained by the technique of [1]. Second, we briefly discuss the tightness of the regret bound in $e \sqrt{n}$ -dependence on time Equation 6 for a general graph with L nodes and E edges. Note that the $O(e)$ is near-optimal, and the $O(d)$ -dependence on feature dimension is standard in linear bandits [1, 33], $? e d$ results are only known for impractical algorithms. The $O(L e ? K)$ factor is due to the since $O(e)$ fact that the reward in this problem is from K to L , rather than from 0 to 1. To explain the $O(E)$ $e e$ factor in this bound, notice that one $O(E)$ factor is due to the fact that at most $O(E)$ edges might be observed at each round (see Theorem 3), and is intrinsic to the problem similarly to combinatorial $e = E$ factor is due to linear generalization (see Lemma 1) and might semi-bandits [19]; another $O(e(d(L ? K)E\sqrt{n}/(??)))$ regret bound in be removed by better analysis. We conjecture that our $O p e = E = d$ away from being tight. this case is at most $O(4.3$

Proof sketch

We now outline the proof of Theorem 1. For each round $t \leq n$, we define the favorable event $T \leq t \leq 1 = \{x \in \{1, \dots, K\} : \exists c \in \mathbb{R}^d \text{ s.t. } M_{t,1} \leq c^T x, \forall c \in \mathbb{R}^d, \forall t \leq t\}$, and the unfavorable event $\bar{T} \leq t \leq 1$ as the complement of $T \leq t \leq 1$. If we decompose $E[R_{t,??}]$, the $(??)$ -scaled expected regret at round t , over events

t_1 and t_1 , and bound R_{t_1} on event t_1 using the naive bound $R_{t_1} \leq L \leq K$, then,

$$E[R_{t_1}] \leq P(t_1) E[R_{t_1} - t_1] + P(t_1) [L \leq K].$$

By choosing c as specified by Equation 4, we have $P(t_1) [L \leq K] \leq 1/n$ (see Lemma 2 in the appendix). On the other hand, notice that by definition of t_1 , $w(e) \leq U_t(e)$, $e \in E$ under event t_1 . Using the monotonicity of f in the probability weight, and the fact that ORACLE is an (ϵ, δ) -approximation algorithm, we have $E[R_{t_1} - t_1] \leq E[f(S_t, U_t) - f(S_t, w) | t_1] / (\epsilon \delta)$. The next observation is that, from the linearity of expectation, the gap $f(S_t, U_t) - f(S_t, w)$ decomposes over nodes $v \in V_{S_t}$. Specifically, for any source node set $S \subseteq V$, any probability weight function $w : E \rightarrow [0, 1]$, and any node $v \in V$, we define $f(S, w, v)$ as the probability that node v is influenced if the source node set is S and the probability weight is w . Hence, we have $P(f(S_t, U_t) - f(S_t, w) | t_1) = \sum_{v \in V_{S_t}} [f(S_t, U_t, v) - f(S_t, w, v)]$. 7

2 12 2 10

8

16 L

24

32

#10 2 11 2.5

2 13

Regret

Regret

Regret

2 14

$\epsilon = 0.7, X = I$

2 15

Star Ray

2 11 29

8

16 L

24

32

(a) Stars and rays: The log-log plots of the n -step regret of IMLinUCB in two graph topologies after $n = 104$ steps. We vary the number of nodes L and the mean edge weight ϵ .

5

$\epsilon = 0.8, X = X_4$ CUCB IMLinUCB with $d=10$

2 102

Cumulative Regret

$\epsilon = 0.8, X = I$

2 16

1.5

29

1

28

8
0.5
16 L
24
32
0 0
1000
2000
3000
4000
5000

Number of Rounds

(b) Subgraph of Facebook network

Figure 2: Experimental results In the appendix, we show that under any weight function, the diffusion process from the source node set S_t to the target node v can be modeled as a Markov chain. Hence, weight function U_t and w give us two Markov chains with the same state space but different transition probabilities. $f(S_t, U_t, v) \leq f(S_t, w, v)$ can be recursively bounded based on the state diagram of the Markov chain under weight function w . With some algebra, Theorem 3 in Appendix A bounds $f(S_t, U_t, v) \leq f(S_t, w, v)$ by the edge-level gap $U_t(e) - w(e)$ on the observed relevant edges for node v , $P f(S_t, U_t, v) \leq P f(S_t, w, v) + \sum_{e \in E_{St,v}} [U_t(e) - w(e)] \mathbb{1}_{\{O_t(e)\}} - H_{t-1}(S_t)$, (9)

for any t , any “history” (past observations) H_{t-1} and S_t such that (9) holds, and any $v \in V_{S_t}$, where $E_{S_t,v}$ is the set of edges relevant to v and $O_t(e)$ is the event that edge e is observed at round t . Based on Equation 9, we can prove Theorem 1 using the standard linear-bandit techniques (see Appendix A).

5

Experiments

In this section, we present a synthetic experiment in order to empirically validate our upper bounds on the regret. Next, we evaluate our algorithm on a real-world Facebook subgraph. 5.1

Stars and rays

In the first experiment, we evaluate IMLinUCB on undirected stars and rays (Figure 1) and validate that the regret grows with the number of nodes L and the maximum observed relevance C as shown in Table 1. We focus on the tabular case ($X = I$) with $K = |S| = 1$, where the IM problem can be solved exactly. We vary the number of nodes L ; and edge weight $w(e) = C$, which is the same for all edges e . We run IMLinUCB for $n = 104$ steps and verify that it converges to the optimal solution in each experiment. We report the n -step regret of IMLinUCB for $8 \leq L \leq 32$ in Figure 2a. Recall that $e \in E_{St,v}$ for star and $R(n) = O(L e^{94})$ for ray. from Table 1, $R(n) = O(L^p)$ We numerically estimate the growth of regret in L , the exponent of L , in the log-log space of L and regret. In particular, since $\log(f(L)) = p \log(L) + \log(c)$ for any $f(L) = cL^p$ and $c \geq 0$, both p and $\log(c)$ can be estimated by linear regression in the new space. For star graphs with $C = 0.8$ and $C = 0.7$, our estimated growth are respectively

$O(L^{2.040})$ and $O(L^{2.056})$, which are close to the e^2). For ray graphs with $\gamma = 0.8$ and $\gamma = 0.7$, our estimated growth are respectively expected $O(L^{2.488})$ and $O(L^{2.467})$, which are again close to the expected $O(L^{\text{observed relevance } C})$ proposed in Section 4.1 is a reasonable complexity metric for these two topologies. 5.2

Subgraph of Facebook network

In the second experiment, we demonstrate the potential performance gain of IMLinUCB in realworld influence maximization semi-bandit problems by exploiting linear generalization across edges. Specifically, we compare IMLinUCB with CUCB in a subgraph of Facebook network from [22]. The subgraph has $L = |V| = 327$ nodes and $|E| = 5038$ directed edges. Since the true probability weight $w(e)$

function w is not available, we independently sample $w(e)$'s from the uniform distribution $U(0, 0.1)$ and treat them as ground-truth. Note that this range of probabilities is guided by empirical evidence in [14, 3]. We set $n = 5000$ and $K = 10$ in this experiment. For IMLinUCB, we choose $d = 10$ and generate edge feature x_e 's as follows: we first use node2vec algorithm [15] to generate a node feature in \mathbb{R}^d for each node $v \in V$; then for each edge e , we generate x_e as the element-wise product of node features of the two nodes connected to e . Note that the linear generalization in this experiment is imperfect in the sense that $\min_{e \in E} |x_e| \leq \max_{e \in E} |x_e| = 0$. For both CUCB and IMLinUCB, we choose ORACLE as the state-of-the-art offline IM algorithm proposed in [27]. To compute the cumulative regret, we compare against a fixed seed set S^* obtained by using the true w as input to the oracle proposed in [27]. We average the empirical cumulative regret over 10 independent runs, and plot the results in Figure 2b. The experimental results show that compared with CUCB, IMLinUCB can significantly reduce the cumulative regret by exploiting linear generalization across $w(e)$'s.

6

Related Work

There exist prior results on IM semi-bandits [21, 8, 31]. First, Lei et al. [21] gave algorithms for the same feedback model as ours. The algorithms are not analyzed and cannot solve large-scale problems because they estimate each edge weight independently. Second, our setting is a special case of stochastic combinatorial semi-bandit with a submodular reward function and stochastically observed edges [8]. Their work is the closest related work. Their gap-dependent and gap-free bounds are both problematic because they depend on the reciprocal of the minimum observation probability p^* of an edge: Consider a line graph with $|E|$ edges where all edge weights are 0.5. Then $1/p^*$ is $2|E|$. On the other hand, our derived regret bounds in Theorem 1 are polynomial in all quantities of interest. A very recent result of Wang and Chen [32] removes the $1/p^*$ factor in [8] for the tabular case and [32] presents a worst-case bound of $O(L|E|n)$, which in the tabular complete graph case improves over our result by $O(L)$. On the other hand, their analysis does not give structural guarantees that we provide with maximum observed relevance C^* obtaining potentially much better results for the case in hand and giving insights for the

complexity of IM bandits. Moreover, both Chen et al. [8] and Wang and Chen [32] do not consider generalization models across edges or nodes, and therefore their proposed algorithms are unlikely to be practical for real-world social networks. In contrast, our proposed algorithm scales to large problems by exploiting linear generalization across edges. IM bandits for different influence models and settings: There exist a number of extensions and related results for IM bandits. We only mention the most related ones (see [28] for a recent survey). Vaswani et al. [31] proposed a learning algorithm for a different and more challenging feedback model, where the learning agent observes influenced nodes but not the edges, but they do not give any guarantees. Carpentier and Valko [5] give a minimax optimal algorithm for IM bandits but only consider a local model of influence with a single source and a cascade of influences never happens. In related networked bandits [11], the learner chooses a node and its reward is the sum of the rewards of the chosen node and its neighborhood. The problem gets more challenging when we allow the influence probabilities to change [2], when we allow the seed set to be chosen adaptively [30], or when we consider a continuous model [12]. Furthermore, Sigla et al. [26] treats the IM setting with an additional observability constraints, where we face a restriction on which nodes we can choose at each round. This setting is also related to the volatile multi-armed bandits where the set of possible arms changes [4]. Vaswani et al. [29] proposed a diffusion-independent algorithm for IM semi-bandits with a wide range of diffusion models, based on the maximum-reachability approximation. Despite its wide applicability, the maximum reachability approximation introduces an additional approximation factor to the scaled regret bounds. As they have discussed, this approximation factor can be large in some cases. Lagr?e et al. [20] treat a persistent extension of IM bandits when some nodes become persistent over the rounds and no longer yield rewards. This work is also a generalization and extension of recent work on cascading bandits [17, 18, 34], since cascading bandits can be viewed as variants of online influence maximization problems with special topologies (chains). Acknowledgements The research presented was supported by French Ministry of Higher Education and Research, Nord-Pas-de-Calais Regional Council and French National Research Agency projects ExTra-Learn (n.ANR-14-CE24-0010-01) and BoB (n.ANR-16-CE23-0003). We would also like to thank Dr. Wei Chen and Mr. Qinshi Wang for pointing out a mistake in an earlier version of this paper.

9

2 References

- [1] Yasin Abbasi-Yadkori, David P?l, and Csaba Szepesv?ri. Improved algorithms for linear stochastic bandits. In *Neural Information Processing Systems*, 2011.
- [2] Yixin Bao, Xiaoke Wang, Zhi Wang, Chuan Wu, and Francis C. M. Lau. Online influence maximization in non-stationary social networks. In *International Symposium on Quality of Service*, apr 2016.
- [3] Nicola Barbieri, Francesco Bonchi, and Giuseppe Manco. Topic-aware social influence propa-

gation models. *Knowledge and information systems*, 37(3):555–584, 2013. [4] Zahy Bnaya, Rami Puzis, Roni Stern, and Ariel Felner. Social network search as a volatile multi-armed bandit problem. *Human Journal*, 2(2):84–98, 2013. [5] Alexandra Carpentier and Michal Valko. Revealing graph bandits for maximizing local influence. In *International Conference on Artificial Intelligence and Statistics*, 2016. [6] Wei Chen, Chi Wang, and Yajun Wang. Scalable influence maximization for prevalent viral marketing in large-scale social networks. In *Knowledge Discovery and Data Mining*, 2010. [7] Wei Chen, Yajun Wang, and Yang Yuan. Combinatorial multi-armed bandit: General framework, results and applications. In *International Conference on Machine Learning*, 2013. [8] Wei Chen, Yajun Wang, and Yang Yuan. Combinatorial multi-armed bandit and its extension to probabilistically triggered arms. *Journal of Machine Learning Research*, 17, 2016. [9] Varsha Dani, Thomas P Hayes, and Sham M Kakade. Stochastic linear optimization under bandit feedback. In *Conference on Learning Theory*, 2008. [10] David Easley and Jon Kleinberg. *Networks, Crowds, and Markets: Reasoning About a Highly Connected World*. Cambridge University Press, 2010. [11] Meng Fang and Dacheng Tao. Networked bandits with disjoint linear payoffs. In *International Conference on Knowledge Discovery and Data Mining*, 2014. [12] Mehrdad Farajtabar, Xiaojing Ye, Sahar Harati, Le Song, and Hongyuan Zha. Multistage campaigning in social networks. In *Neural Information Processing Systems*, 2016. [13] M Gomez Rodriguez, B Schölkopf, Langford J Pineau, et al. Influence maximization in continuous time diffusion networks. In *International Conference on Machine Learning*, 2012. [14] Amit Goyal, Francesco Bonchi, and Laks VS Lakshmanan. Learning influence probabilities in social networks. In *Proceedings of the third ACM international conference on Web search and data mining*, pages 241–250. ACM, 2010. [15] Aditya Grover and Jure Leskovec. node2vec: Scalable feature learning for networks. In *Knowledge Discovery and Data Mining*. ACM, 2016. [16] David Kempe, Jon Kleinberg, and Eva Tardos. Maximizing the spread of influence through a social network. *Knowledge Discovery and Data Mining*, page 137, 2003. [17] Branislav Kveton, Csaba Szepesvari, Zheng Wen, and Azin Ashkan. Cascading bandits: Learning to rank in the cascade model. In *Proceedings of the 32nd International Conference on Machine Learning*, 2015. [18] Branislav Kveton, Zheng Wen, Azin Ashkan, and Csaba Szepesvari. Combinatorial cascading bandits. In *Advances in Neural Information Processing Systems* 28, pages 1450–1458, 2015. [19] Branislav Kveton, Zheng Wen, Azin Ashkan, and Csaba Szepesvari. Tight regret bounds for stochastic combinatorial semi-bandits. In *Proceedings of the 18th International Conference on Artificial Intelligence and Statistics*, 2015. [20] Paul Lagrèe, Olivier Cappé, Bogdan Cautis, and Silviu Maniu. Effective large-scale online influence maximization. In *International Conference on Data Mining*, 2017. 10 [21] Siyu Lei, Silviu Maniu, Luyi Mo, Reynold Cheng, and Pierre Senellart. Online influence maximization. In *Knowledge Discovery and Data mining*, 2015. [22] Jure Leskovec and Andrej Krevl. Snap datasets: Stanford large network dataset collection. <http://snap.stanford.edu/data>, jun 2014. [23] Yanhua Li, Wei Chen, Yajun Wang, and Zhi-Li Zhang. Influence diffusion dynamics

and influence maximization in social networks with friend and foe relationships. In ACM international conference on Web search and data mining. ACM, 2013. [24] Praneeth Netrapalli and Sujay Sanghavi. Learning the graph of epidemic cascades. In ACM SIGMETRICS Performance Evaluation Review, volume 40, pages 211?222. ACM, 2012. [25] Kazumi Saito, Ryohei Nakano, and Masahiro Kimura. Prediction of information diffusion probabilities for independent cascade model. In Knowledge-Based Intelligent Information and Engineering Systems, pages 67?75, 2008. [26] Adish Singla, Eric Horvitz, Pushmeet Kohli, Ryan White, and Andreas Krause. Information gathering in networks via active exploration. In International Joint Conferences on Artificial Intelligence, 2015. [27] Youze Tang, Xiaokui Xiao, and Shi Yanchen. Influence maximization: Near-optimal time complexity meets practical efficiency. 2014. [28] Michal Valko. Bandits on graphs and structures. habilitation, ?cole normale sup?rieure de Cachan, 2016. [29] Sharan Vaswani, Branislav Kveton, Zheng Wen, Mohammad Ghavamzadeh, Laks VS Lakshmanan, and Mark Schmidt. Model-independent online learning for influence maximization. In International Conference on Machine Learning, 2017. [30] Sharan Vaswani and Laks V. S. Lakshmanan. Adaptive influence maximization in social networks: Why commit when you can adapt? Technical report, 2016. [31] Sharan Vaswani, Laks. V. S. Lakshmanan, and Mark Schmidt. Influence maximization with bandits. In NIPS workshop on Networks in the Social and Information Sciences 2015, 2015. [32] Qinshi Wang and Wei Chen. Improving regret bounds for combinatorial semi-bandits with probabilistically triggered arms and its applications. In Neural Information Processing Systems, mar 2017. [33] Zheng Wen, Branislav Kveton, and Azin Ashkan. Efficient learning in large-scale combinatorial semi-bandits. In International Conference on Machine Learning, 2015. [34] Shi Zong, Hao Ni, Kenny Sung, Nan Rosemary Ke, Zheng Wen, and Branislav Kveton. Cascading bandits for large-scale recommendation problems. In Uncertainty in Artificial Intelligence, 2016.